

# Assignment 4: The Ocean

## Practical information

Deadline: Saturday 9/3 23.59

Resources:

- ERDA for file storage
- Nvidia profiler to determine the parallelisation bottlenecks
- Jupyter for the Terminal to access DAG Nvidia vGPU instances

Handin:

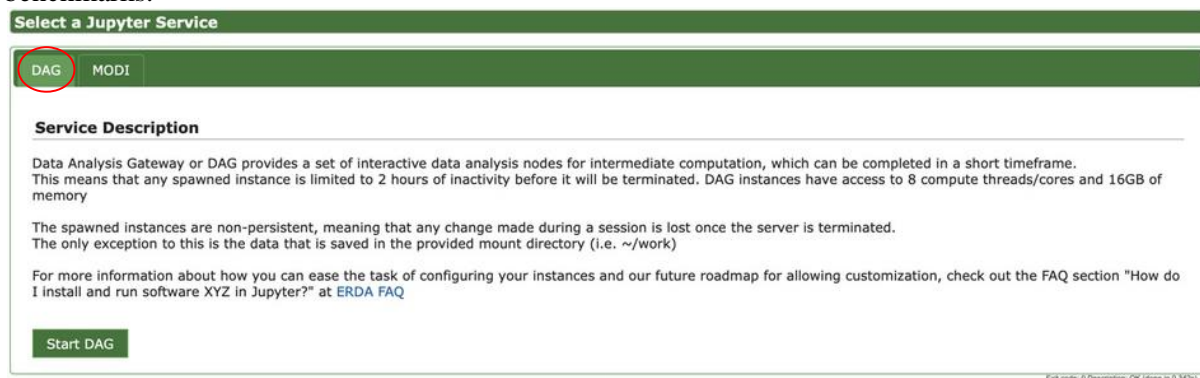
- Total assignment: a report of up to 3 pages in length (excluding the code)
- Use the template on Absalon to include your code in the report

## Introduction

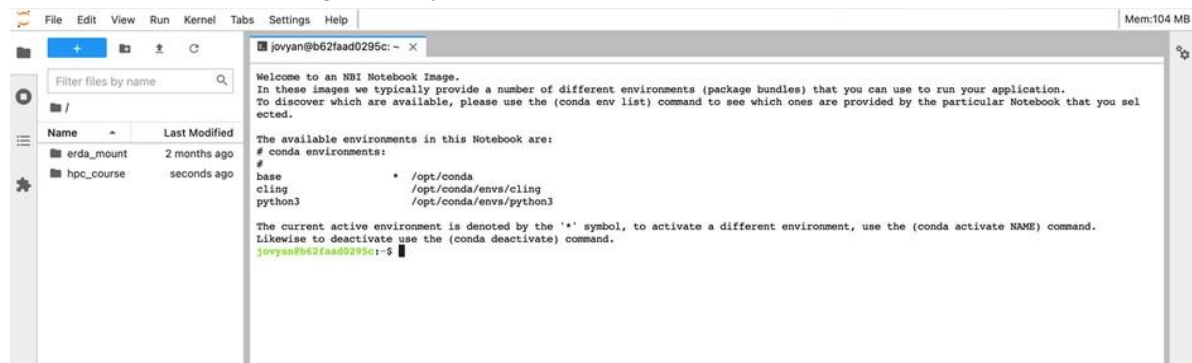
The Shallow Water (SW) model (section 13.3) is the simplest numerical representation of the ocean. Still, it has reasonable skills when used to predict the evolution of storm surges or Tsunamis. Moreover, it illustrates nicely the functioning and parallelization of stencil operations.

## DAG

For this assignment we need `nvcc++` to compile and DAG for source code profiling and running the benchmarks.



You can read more about DAG in the user guide: <https://erda.dk/public/ucph-erda-user-guide.pdf>. Spin up a Jupyter session on DAG selecting the “HPC GPU notebook” notebook image. In the terminal (or the folder view on the right side) you can see a number of folders.



The different folders contain:

- `erda_mount`: your own files.
- `hpc_course`: course files.

## Preparations

Start by copying the exercise to your storage area and enter in to the folder. You can write 'ls' to get a file listing of the folder.

```
Makefile
sw_parallel.cpp
sw_sequential.cpp
visualize.ipynb
```

Before you can run the code, you need to compile it. This can be done by running `make` in the terminal. The `sw_sequential.cpp` code is identical to `sw_parallel.cpp` there (with produced corresponding binaries `sw_sequential` and `parallel`) to give you a backup. The `visualize.ipynb` is for SW model output visualisation and analysis.

To run the code for 500 time-steps on DAG and write the model output in ASCII file to your storage you can do:

```
./sw_sequential --iter 500 --out sw_output.txt
```

## Nvidia profiler

NVIDIA profiler enables you to understand and optimize the performance of your OpenACC application. An example of command-line `nvprof` profiler output for parallelised SW model is given below. One can, for example, see a runtime of four compute kernels

- `integrate_116_gpu`
- `integrate_123_gpu`
- `exchange_vertical_ghost_lines_100_gpu`
- `exchange_horizontal_ghost_lines_87_gpu`

and time spent for copying data from Host to Device and back, lines with `[CUDA memcpy HtoD]` and `[CUDA memcpy DtoH]`, respectively. *Please note that `nvprof` does not work anymore. Check the last slide of lecture 10 for `nsys` instead.*

```
jovyan@b62fadd0295c:~/erda_mount/Teaching/hpc_course_private/module5$ nvprof ./sw_parallel --iter 500
==534== NVPROF is profiling process 534, command: ./sw_parallel --iter 500
checksum: 4117.75
elapsed time: 1.06693 sec
==534== Profiling application: ./sw_parallel --iter 500
==534== Profiling result:
Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities:  43.05%  16.155ms      500  32.310us  32.032us  33.183us  integrate_116_gpu(Water&)
                34.99%  13.129ms      500  26.258us  25.984us  26.880us  integrate_123_gpu(Water&)
                10.59%  3.9730ms     1000  3.9730us  3.9030us  14.720us  exchange_vertical_ghost_lines_100_gpu(std::array<std::array<float,
unsigned long=512>, unsigned long=512>*)
                10.02%  3.7595ms     1000  3.7590us  3.7110us  5.6640us  exchange_horizontal_ghost_lines_87_gpu(std::array<std::array<float
, unsigned long=512>, unsigned long=512>*)
                0.68%  254.56us        1  254.56us  254.56us  254.56us  [CUDA memcpy HtoD]
                0.67%  252.19us        1  252.19us  252.19us  252.19us  [CUDA memcpy DtoH]
API calls:      69.07%  203.85ms      7002  7.6510us  1.0040us  635.72us  cuDevicePrimaryCtxRetain
                7.98%  23.540ms        1  23.540ms  23.540ms  23.540ms  cuStreamSynchronize
                4.36%  12.860ms      3000  4.2860us  3.4370us  533.67us  cuMemHostAlloc
                0.28%  838.18us        1  838.18us  838.18us  838.18us  cuLaunchKernel
                0.08%  225.78us        2  112.89us  109.45us  116.33us  cuMemAllocHost
                0.05%  148.92us        1  148.92us  148.92us  148.92us  cuMemAlloc
                0.01%  30.053us        1  30.053us  30.053us  30.053us  cuModuleLoadDataEx
                0.00%  12.201us        3  4.0670us  2.1660us  7.5830us  cuMemcpyHtoDAsync
                0.00%  10.629us        1  10.629us  10.629us  10.629us  cuEventRecord
                0.00%  8.4790us        2  4.2390us  1.8640us  6.6150us  cuMemcpyDtoHAsync
                0.00%  7.6780us        1  7.6780us  7.6780us  7.6780us  cuDeviceGetPCIBusId
                0.00%  6.4870us        3  2.1620us  826ns    3.1350us  cuPointerGetAttributes
                0.00%  4.7920us        4  1.1980us  392ns    2.9810us  cuEventCreate
                0.00%  4.2280us       10  422ns    156ns    2.1590us  cuModuleGetFunction
                0.00%  3.0310us        4  757ns    149ns    2.5130us  cuDeviceGetAttribute
                0.00%  2.6640us        3  888ns    211ns    1.8230us  cuDeviceGet
                0.00%  2.2560us        1  2.2560us  2.2560us  2.1050us  cuDeviceGetCount
                0.00%  2.1050us        1  2.1050us  2.1050us  2.1050us  cuEventSynchronize
                0.00%  2.0000us        3  666ns    299ns    908ns    cuCtxGetCurrent
                0.00%  526ns          2  263ns    175ns    351ns    cuCtxSetCurrent
                0.00%  231ns          1  231ns    231ns    231ns    cuDeviceComputeCapability
                0.00%  24.159ms        1  24.159ms  24.159ms  24.159ms  cuDriverGetVersion
OpenACC (excl):  19.91%  24.159ms        1  24.159ms  24.159ms  24.159ms  acc_enter_data@sw_parallel.cpp:145
                16.83%  20.430ms     1500  13.619us  1.7520us  57.933us  acc_wait@sw_parallel.cpp:116
                15.25%  18.508ms     1500  12.338us  1.7480us  637.24us  acc_wait@sw_parallel.cpp:123
                8.48%  10.290ms     2000  5.1440us  1.8420us  25.397us  acc_wait@sw_parallel.cpp:100
                8.45%  10.256ms     2000  5.1280us  1.8510us  25.614us  acc_wait@sw_parallel.cpp:87
                4.91%  5.9630ms     1000  5.9630us  4.6990us  535.39us  acc_enqueue_launch@sw_parallel.cpp:87 (_Z38exchange_horizontal_gho
st_lines_87_gpuRSt5arrayIS_ifLm512EE)
                4.37%  5.3077ms     1000  5.3070us  4.6310us  23.830us  acc_enqueue_launch@sw_parallel.cpp:100 (_Z37exchange_vertical_ghos
t_lines_100_gpuRSt5arrayIS_ifLm512EE)
                2.71%  3.2845ms        500  6.5680us  4.8360us  592.57us  acc_enqueue_launch@sw_parallel.cpp:116 (_Z17integrate_116_gpuR5Wat
```

## Task 1: OpenACC parallelise the program (points 5)

The key challenge is to identify which parts of the code can reasonably be executed by the GPUs and to find suitable OpenACC directives and clauses for optimal parallelization. With the help of a profiler determine the bottlenecks. Play around a bit with the `#pragma` and see if you can improve on your first try. Thus, you need to save the profiler output of your various experiments.

## Task 2: Strong and weak scaling using SLURM (points 5)

Measure the weak and strong scaling of your programs. Change the variables `Nx` and `Ny` to measure weak scaling, you should think about and explain your choice of grid size and how this map to the available vGPUs. For strong scaling you can use the clauses `num_gangs()` and `vector()` to control the number of thread blocks and sizes of thread blocks. Some key figures to note: You have 14 streaming multiprocessors (compute units or SMs) available, the maximum number of threads per thread-block is 1024 and each multiprocessor can handle at most 2048 threads.