

High Performance Parallel Computing 2023

Introduction to C++ and Tools

Markus Jochum & Troels Haugbølle (slides courtesy Rasmus Munk)

Niels Bohr Institute
University of Copenhagen

Copenhagen, February 6, 2023

Practical Information

- Teachers

Troels Haugbølle (haugbel@nbi.ku.dk)

Markus Jochum (mjochum@nbi.ku.dk)

Jorge Expósito Patiño (jepa@di.ku.dk)

Roman Nuterman (nuterman@nbi.ku.dk)

- Location

see Absalon - Course Overview

- Lectures

Monday 13:15 - 16:00

Wednesday 10:15 - 12:00

- Exercises

Wednesday 13:15 - 16:00

Modules

- Week 1: Basic architecture and programming (C++11)
- Week 2: Vectorization – architecture and programming (OpenMP SIMD)
- Week 3: Task Farming (MPI Master-Worker)
- Week 4: Shared memory architectures and programming (OpenMP)
- Week 5: GPUs and many-Core architecture (OpenACC)
- Week 6: Distributed Memory and Networked Architectures (MPI)
- Week 7: Project Week

Examination

- Requirement
 - 6 assignments, reports in groups of up to 3 persons (max three pages) 6 x 10%
 - 1 week-long project in groups of 4 to 6 persons (max 10 pages) 40%
- Assignment 1: Epidemic Model SIR (C++ programming)
Deadline 12/2
- Assignment 2: Molecular Dynamics (Vectorization and Memory layout)
Deadline 19/2
- Assignment 3: Particle Physics Electron data (Task Farming with MPI Master-Worker)
Deadline 26/2
- Assignment 4: Seismology (Shared Memory parallelisation with OpenMP)

Deadline 5/3
- Assignment 5: Shallow Water (GPU acceleration with OpenACC)
Deadline 12/3
- Assignment 6: Climate Model (Distributed Memory with MPI)
Deadline 19/3

Today's Lecture

- Introduction of C++11
- Makefile
- Debugging
- Jupyter – use the cloud

What is High Performance Parallel Computing?



- Solving tightly coupled problems using tightly coupled systems
- Handle domain decomposition
- Handle communication latency
- Handle large data sets

What is C++?

- A general-purpose object-oriented programming language.
- Invented by Bjarne Stroustrup (1979).
- Extension of the C language (but not a superset).
- Initially, the language was called “C with Classes”.
- Encapsulates both high- and low-level language features.

Why C++?

Pros:

- Link compatible with C and FORTRAN
- Performance similar to C and FORTRAN
- Great Standard Library
- Static typed
- Predictable performance e.g. no garbage collector
- High-level containers
- Generic programming through templates
- Object-oriented

Cons:

- Object-oriented
- HUGE and Complex!

A Very Short Introduction to C++

- We will only cover a fraction of C++
- `en.cppreference.com`
- `www.cplusplus.com`
- `erlerobotics.gitbooks.io/erle-robotics-cpp-gitbook`
- `www3.ntu.edu.sg/home/ehchua/programming/index.html`

Python versus C++

```
def abs_add(a, b):  
    result = a + b  
    if result < 0:  
        result = -result  
  
    return result
```

```
int abs_add(int a, int b) {  
    int result = a + b;  
    if(result < 0) {  
        result = -result;  
    }  
    return result;  
}
```

- Newline versus semicolons
- Tabs versus brackets
- Dynamic typed versus static typed

Hello World Example

```
#include <iostream>

int main(int argc, char **argv) {
    std::cout << "Hello World!\n";
    return 0;
}
```

- The execution starts at `main()`
- Use `#include <...>` to include libraries
- Use `std::cout` to write to standard out
- Use `"\n"` to write newline
- `int argc, char **argv` are the command line arguments

```
#include <stdint>
#include <complex>

char;           // 1 byte
short;          // 2 bytes (compiler specific)
int; long int;  // 4 bytes (compiler specific)
long long int;  // 8 bytes (compiler specific)
float;          // 4 bytes
double;         // 8 bytes

// Unsigned versions
unsigned char; unsigned short; unsigned int...

// C99 types from <stdint>
int8_t; int16_t; int32_t; int64_t;
uint8_t; uint16_t; uint32_t; uint64_t;

// Complex types from <complex>
std::complex<float> // 8 bytes
std::complex<double> // 16 bytes
```

References

```
void add_inplace(int &a, int b) {  
    a = a + b;  
}  
  
int main() {  
    int v = 0; // Declare a new integer named `v`  
    add_inplace(v, 42); // Updating the value of `v`  
    std::cout << "The new value of `v` is: " << v << "\n";  
}
```

- `int v = 0;` both declare and initiate a new integer.
- `add_inplace` doesn't return anything, instead it updates value in-place.
- The type `int &` is a *reference* to an integer.

Constant Reference

```
void add_inplace(int &a, const int &b) {  
    a = a + b;  
}  
  
int main() {  
    int v = 0; // Declare a new integer named `v`  
    add_inplace(v, 42); // Updating the value of `v`  
    std::cout << "The new value of `v` is: " << v << "\n";  
}
```

The type `const int &` is a *constant* reference to an integer.

```
int main() {  
    int v = 0;  
    int &ref = v;  
    const int &cref = v;  
}
```

Pointers

```
void add_inplace(int *a, const int b) {  
    *a = *a + b;  
}  
  
int main() {  
    int v = 0;  
    add_inplace(&v, 42); // Updating the value of `v`  
    std::cout << "The new value of `v` is: " << v << "\n";  
}
```

- The type `int *` is a *pointer* to an integer.
- The `&v` operator provides the *pointer* to 'v'

For Loop

```
int s = 0;
for(int i=0; i<10; ++i) {
    s += i;
}
```

```
s = 0
for i in range(10):
    s += i
```

While Loop

```
int s = 0;
int i = 0;
while(i < 10) {
    s += i;
    ++i;
}
```

```
s = 0
i = 0
while i < 10:
    s += i
    i += 1
```



```
#include <vector>    // Standard Vector Class

// Empty vector of integers: []
std::vector<int> vec;

// Uninitiated vector of integers of size three
std::vector<int> vec(3);
std::vector<int> vec{3};

// Initiated vector of integers of size three: [42, 42, 42]
std::vector<int> vec(3, 42);

// Initiated vector of integers of size three: [2, 4, 3]
std::vector<int> vec = {2, 4, 3};

// Return the size of the vector
vec.size()

// Append a value to the vector
vec.push_back(42)
```

For-each-loop

For each element `v` in `vec`, run the loop body.

```
void f(std::vector vec) {  
    int s = 0;  
    for(int &v: vec) {  
        s += v;  
    }  
}
```

```
def f(vec):  
    s = 0  
    for v in vec:  
        s += v
```

Question

```
int vec_sum(std::vector vec) {  
    int s = 0;  
    for(int &v: vec) {  
        s += v;  
    }  
    return s;  
}
```

How can we optimize the performance of `vec_sum`?

Answer - use a reference

```
int vec_sum(std::vector &vec) {  
    int s = 0;  
    for(int &v: vec) {  
        s += v;  
    }  
    return s;  
}
```

We can avoid copying `vec` by using a reference

```
std::vector &
```

Function Overload

```
int add_one(int a) {  
    return a + 1;  
}  
  
float add_one(float a) {  
    return a + 1;  
}  
  
// Print `43`  
std::cout << add_one(42) << "\n";  
// Print `5.2`  
std::cout << add_one(4.2) << "\n";
```

Generic Programing – Template

```
template<typename T>
T add_one(T a) {
    return a + 1;
}
// Print `43`
std::cout << add_one(42) << "\n";
// Print `5.2`
std::cout << add_one(4.2) << "\n";
// Print `5`
std::cout << add_one<int>(4.2) << "\n";
```

`std::vector<int>` is a templated class.

Class

```
class Body {  
public:  
    double mass;  
    double pos_x;  
    double pos_y;  
    double vel_x;  
    double vel_y;  
};  
// Uninitiated instantiation of `Body`  
Body my_body;  
// Initiated instantiation of `Body`  
Body my_body{42, 2.4, 1.4, 0, 0};  
// Write to the `pos_x` attribute of `my_body`  
my_body.pos_x = 4.2;
```

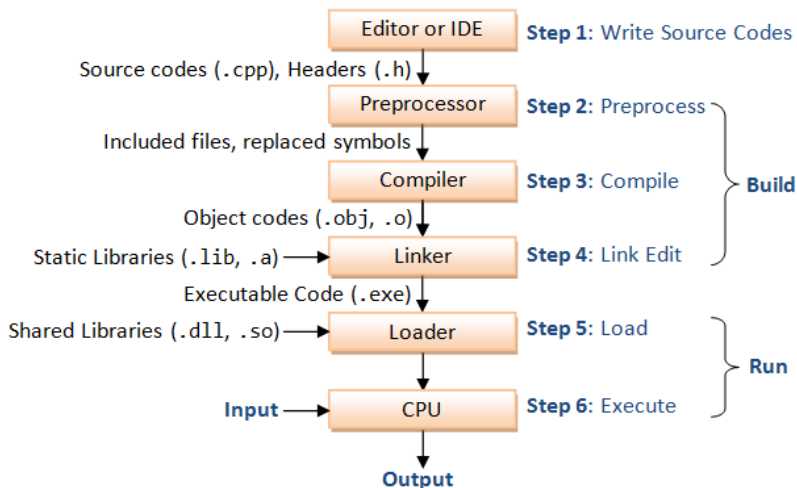
Class Constructor

```
class Account {  
private:  
    double balance;  
public:  
    Account(double start_balance) : balance{start_balance} {  
        std::cout << "Start balance is: " << balance << "\n";  
    }  
};  
  
// Create a new account with a balance of 4.2  
Account my_account{4.2};  
my_account.balance; // ERROR!
```


Class Methods

```
class Account {  
private:  
    double balance;  
public:  
    Account(double start_balance) : balance{start_balance} {  
        std::cout << "Start balance is: " << balance << "\n";  
    }  
  
    void deposit(double amount) {  
        balance += amount;  
    }  
  
    void withdraw(double amount) {  
        balance -= amount;  
    }  
};  
Account my_account{4.2};  
my_account.withdraw(2);
```

The Compiling Process



Compilation

```
#include <iostream>

int main(int argc, char **argv) {
    std::cout << "Hello World!\n";
    return 0;
}
```

```
g++ -std=c++11 -O3 -g hello.cpp -o hello
```

- GNU Compiler Collection `gcc`. Alternatives: `clang`, `icc`, `pgcc`
- Optimization flag `-O0`, `-O1`, `-O2`, `-O3`
- Debug symbols `-g`

```
LIB ?= -lm
INC ?= -I/opt/hpc_course/include
FLAGS ?= -O3 -DNDEBUG -g -Wall

all: nbody_seq nbody_acc

nbody_seq: nbody_seq.cpp
    pgc++ $(FLAGS) $(LIB) $(INC) -o nbody_seq nbody_seq.cpp

nbody_acc: nbody_acc.cpp
    pgc++ -acc -ta=multicore \
        $(FLAGS) $(LIB) $(INC) -o nbody_acc nbody_acc.cpp

clean:
    rm -f nbody_seq nbody_acc *.o

.PHONY: clean all
```

make

Printing Debugging

Follow the state of your program by using `std::cout`.

```
#include <iostream>

int abs_add(int a, int b) {
    int result = a + b;
    if(result < 0) {
        std::cout << "result is negative: " << result << "\n";
        result = result;
    }
    std::cout << "result of abs_add(): " << result << "\n";
    return result;
}
```

Assertions

Insert `assert()`

```
#include <cassert>

int abs_add(int a, int b) {
    int result = a + b;
    if(result < 0) {
        result = result;
    }
    assert(result >= 0);
    return result;
}
```

NOT setting `-DNDEBUG` enables `assert()`. Use `-Wall` to get all compiler warnings.

```
gcc -DNDEBUG -Wall
```

The GNU Project Debugger (GDB)

```
gdb --args /my_program arg1 arg2 arg3
```

`run` run your program.

`step` run to the next operation.

`next` run to the next line.

`break` set breakpoint.

`list` show current source code.

`bt` show backtrace.

`p` run command and show the result e.g. `p my_var` will show the value of `my_var`.

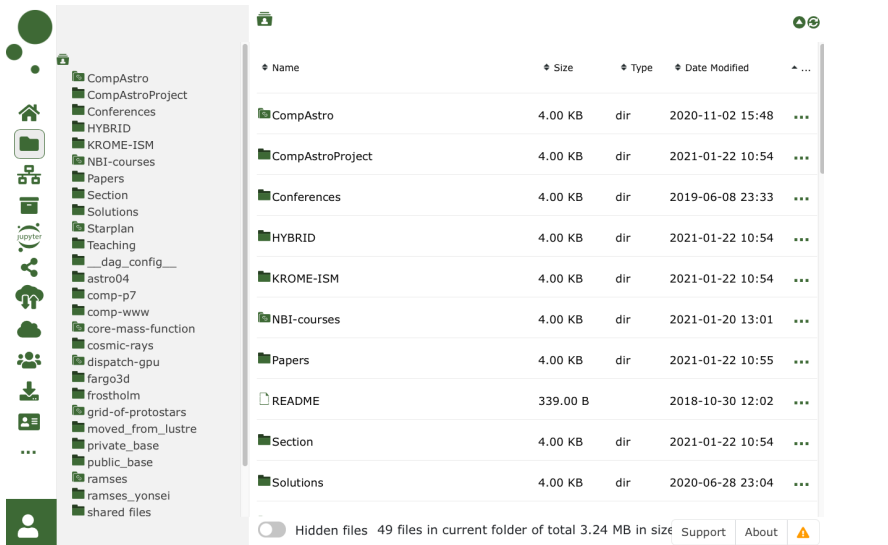
Valgrind

Check your program for illegal memory accesses and memory leaks

```
valgrind /my_program arg1 arg2 arg3
```


```
==21366== HEAP SUMMARY:
==21366==      in use at exit: 26,935 bytes in 31 blocks
==21366==    total heap usage: 535 allocs, 504 frees, 98,010 bytes allocated
==21366==
==21366== LEAK SUMMARY:
==21366==    definitely lost: 0 bytes in 0 blocks
==21366==    indirectly lost: 0 bytes in 0 blocks
==21366==    possibly lost: 0 bytes in 0 blocks
==21366==    still reachable: 26,935 bytes in 31 blocks
==21366==          suppressed: 0 bytes in 0 blocks
==21366== Rerun with --leak-check=full to see details of leaked memory
==21366==
==21366== For counts of detected and suppressed errors, rerun with: -v
==21366== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```


Jupyter



The JupyterLab interface displays a file browser on the left and a file list on the right. The file browser shows a tree view of the file system, including folders like CompAstro, CompAstroProject, Conferences, HYBRID, KROME-ISM, NBI-courses, Papers, Section, Solutions, Starplan, Teaching, __dag_config__, astro04, comp-p7, comp-www, core-mass-function, cosmic-rays, dispatch-gpu, fargo3d, frosthalm, grid-of-protostars, moved_from_lustre, private_base, public_base, ramses, ramses_yonsei, and shared files. The file list on the right shows the contents of the current folder, including folders like CompAstro, CompAstroProject, Conferences, HYBRID, KROME-ISM, NBI-courses, Papers, and Solutions, and a file named README. The file list has columns for Name, Size, Type, Date Modified, and a menu icon. The README file is highlighted.

Name	Size	Type	Date Modified	...
CompAstro	4.00 KB	dir	2020-11-02 15:48	...
CompAstroProject	4.00 KB	dir	2021-01-22 10:54	...
Conferences	4.00 KB	dir	2019-06-08 23:33	...
HYBRID	4.00 KB	dir	2021-01-22 10:54	...
KROME-ISM	4.00 KB	dir	2021-01-22 10:54	...
NBI-courses	4.00 KB	dir	2021-01-20 13:01	...
Papers	4.00 KB	dir	2021-01-22 10:55	...
README	339.00 B		2018-10-30 12:02	...
Section	4.00 KB	dir	2021-01-22 10:54	...
Solutions	4.00 KB	dir	2020-06-28 23:04	...

Hidden files 49 files in current folder of total 3.24 MB in size [Support](#) [About](#) 

Select a Jupyter Service

DAG

MODI

Service Description

Data Analysis Gateway or DAG provides a set of interactive data analysis nodes for intermediate computation, which can be completed in a short timeframe.


This means that any spawned instance is limited to 2 hours of inactivity before it will be terminated. DAG instances have access to 8 compute cores and 8GB of memory

The spawned instances are non-persistent, meaning that any change made during a session is lost once the server is terminated.

The only exception to this is the data that is saved in the provided mount directory (i.e. ~/work)

For more information about how you can ease the task of configuring your instances and our future roadmap for allowing customization, check out the FAQ section "How do I install and run software XYZ in Jupyter?" at [ERDA FAQ](#)

Start DAG



Home

Token

Logout

Spawner Options

Select a notebook image:

HPC Notebook

Spawn

File Edit View Run Kernel Tabs Settings Help

Launcher

cxj_exercises_i.ipynb cxj_exercises_ii.ipynb

Notebook

- Python 3
- C++11
- C++14
- C++17

Console

- Python 3
- C++11
- C++14
- C++17

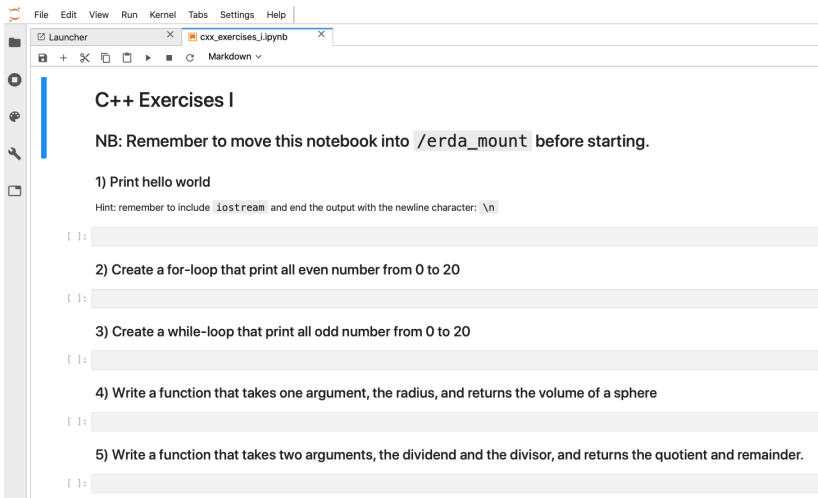
Other

- Terminal
- Text File (Create a new text file)
- Markdown File
- Show Contextual Help

- **NB:** remember to move your files into `erda_mount`

The screenshot shows the JupyterLab interface. On the left is a file explorer showing the directory structure: `/... / module1 / exercises /`. Two files, `cxx_exercises_ii.ipynb` and `cxx_exercises_i.ipynb`, are listed, both modified 6 hours ago. A right-click context menu is open over `cxx_exercises_i.ipynb`, showing options like Open, Open With, Open in New Browser Tab, Rename, Delete, Cut, Copy, Duplicate, Download, Shut Down Kernel, Copy Shareable Link, Copy Path, Copy Download Link, New Folder, and Paste. The main workspace on the right is titled `hpc_course/module1/exercises` and contains three panels: **Launcher**, **Console**, and **Other**. The **Launcher** panel shows a **Notebook** section with four icons for Python 3, C++11, C++14, and C++17. The **Console** panel also shows icons for Python 3, C++11, C++14, and C++17. The **Other** panel shows icons for Terminal, Text File, Markdown File, and Show Contextual Help.

● **NB:** move `cxx_exercises_X.ipynb` into your `erda_mount` !



The screenshot shows a Jupyter Notebook interface. The top bar contains a menu (File, Edit, View, Run, Kernel, Tabs, Settings, Help) and a tab for 'cxz_exercises_i.ipynb'. The left sidebar has icons for file operations and a search icon. The notebook content is as follows:

C++ Exercises I

NB: Remember to move this notebook into `/erda_mount` before starting.

1) Print hello world

Hint: remember to include `iostream` and end the output with the newline character: `\n`

```
[ ]:
```

2) Create a for-loop that print all even number from 0 to 20

```
[ ]:
```

3) Create a while-loop that print all odd number from 0 to 20

```
[ ]:
```

4) Write a function that takes one argument, the radius, and returns the volume of a sphere

```
[ ]:
```

5) Write a function that takes two arguments, the dividend and the divisor, and returns the quotient and remainder.

```
[ ]:
```

At the bottom right, there are navigation icons for the notebook.